

# Scikit-learn Intro

We want to get a machine to learn how to recognize digits.

Scikit-learn has some builtin datasets

```
In [1]: from sklearn import datasets
```

The most famous one in ML is the MNIST dataset.

```
In [2]: digits = datasets.load_digits()
```

Let's take a look at the data briefly.

```
In [3]: print(digits.data)
[[ 0.  0.  5. ...  0.  0.  0.]
 [ 0.  0.  0. ... 10.  0.  0.]
 [ 0.  0.  0. ... 16.  9.  0.]
 ...
 [ 0.  0.  1. ...  6.  0.  0.]
 [ 0.  0.  2. ... 12.  0.  0.]
 [ 0.  0. 10. ... 12.  1.  0.]]
```

That was not very helpful. We need to study it in a bit more detail.

```
In [4]: digits.data.shape
```

```
Out[4]: (1797, 64)
```

```
In [5]: digits.data[0]
```

```
Out[5]: array([ 0.,  0.,  5., 13.,  9.,  1.,  0.,  0.,  0.,  0., 13., 15., 10.,
                15.,  5.,  0.,  0.,  3., 15.,  2.,  0., 11.,  8.,  0.,  0.,  4.,
                12.,  0.,  0.,  8.,  8.,  0.,  0.,  5.,  8.,  0.,  0.,  9.,  8.,
                 0.,  0.,  4., 11.,  0.,  1., 12.,  7.,  0.,  0.,  2., 14.,  5.,
                10., 12.,  0.,  0.,  0.,  0.,  6., 13., 10.,  0.,  0.,  0.]
```

And what are we trying to predict?

```
In [6]: digits.target
```

```
Out[6]: array([0, 1, 2, ..., 8, 9, 8])
```

```
In [7]: digits.target[-10:]
```

```
Out[7]: array([5, 4, 8, 8, 4, 9, 0, 8, 9, 8])
```

```
In [8]: digits.target.shape
```

```
Out[8]: (1797,)
```

## Learning the digits

Let's load a builtin classifier -- an object that decides what the image corresponds to.

```
In [9]: from sklearn import svm
```

Don't worry about the exact values we've put in for the parameters. They're not important for us at the moment.

```
In [10]: clf = svm.SVC(gamma=0.001, C=100.)
```

Let's tell the classifier to learn from the data. We'll show it all the images and tell it which ones they are -- except for the last one. We'll hold that secret.

```
In [11]: clf.fit(digits.data[:-1], digits.target[:-1])
```

```
Out[11]: SVC(C=100.0, cache_size=200, class_weight=None, coef0=0.0,
  decision_function_shape='ovr', degree=3, gamma=0.001, kernel='rbf',
  max_iter=-1, probability=False, random_state=None, shrinking=True,
  tol=0.001, verbose=False)
```

The machine has learned the digits -- or so it thinks. Let's test it on the one digits we never showed it. Notice that we're not telling it the target.

```
In [12]: clf.predict(digits.data[-1:])
```


```
Out[12]: array([8])
```

Let's see what that image looks like and whether the prediction makes sense.

```
In [13]: %matplotlib inline
import numpy as np
from matplotlib import pyplot as plt

plt.figure(figsize=(2, 2))
plt.imshow(digits.images[-1], interpolation='nearest', cmap=plt.cm.binary)
```

```
Out[13]: <matplotlib.image.AxesImage at 0x7f85f98219b0>
```



Let's go back and see where all the learning happened.